


Evidence-grounded Financial Operations Intelligence Agent

GOAI 2026 · Boundless Agents · AI+ 金融

Question → Data → Analysis → Evidence → Decision



Not chat with a CSV.
An auditable investigation.

The gap: dashboards show what, analysts still investigate why

01 Find tables

02 Define KPI

03 Compare periods

04 Drill dimensions

05 Test significance

06 Inspect events

07 Reconcile numbers

08 Write report

A generic LLM can generate an explanation faster — but plausibility is not auditability.

Existing stack

**Dashboard + SQL +
BI**

Missing layer

**Reproducible
investigation**

Design target

**Evidence before
narrative**

FitzSight turns a business question into a verified investigation



Planner decides what to investigate. Deterministic tools own every number. Verifier decides what may be shown.

Planner authority

Approved actions only
No write from SQL or high-impact actions

Calculation authority

Read-only SQL / Python
Statistics and decompositions are executable

Answer authority

EvidenceClaimVerifier
Verification failure → answer withheld

Demo 1 — Why did European FTD conversion deteriorate after July 15?



Investigation evidence

- Response latency shifted by +29.15 minutes after the routing change
- Team A is the largest negative team-level contributor
- Nearby CRM routing event exists in the business-event log
- 8 post-change response-time days exceed the robust threshold

Result: supported root-cause candidate, with evidence and a causal boundary.

Demo 2 — Why did European net deposits fall in the week starting August 3?

Net-deposit change

-\$187.8k

Deposits

+\$59.2k

Withdrawals

+\$246.9k

Top 11 withdrawal share

91.6%

Concentration, not motive inference

Driver decomposition

Deposit change

+59.2k

Withdrawal pressure

-246.9k

Net change

-187.8k

What FitzSight refuses to infer

Customer motive

AML suspicion

Investment intent

Automated account action

Observed driver: concentrated withdrawal pressure. Unsupported story: why customers withdrew.

Demo 3 — Customer Intelligence without hidden labels

Europe customers

6,770

Coverage

100%

High Value users

3.7%

High Value deposit share

53.7%

Transparent behavioral value policy

Deposit value

Deposit frequency

Trade volume

Trade frequency

Withdrawal value

Descriptive use only

- No credit decision
- No AML decision
- No eligibility decision
- No adverse action

Normal Agent SQL never reads customer_segment_gt. Segments are built from observable behavior.

Demo 4 — More leads ≠ better acquisition quality

Lead volume

+315%

Aggregate FTD

-10.84 pp

Paid Search mix

+60.52 pp

Paid Search FTD

-16.44 pp

p-value

4.43e-05

FitzSight separates three questions

1. Volume

2. Channel mix

3. Within-channel
performance

Paid Search is the measurable quality failure — not simply a high lead count.

Demo 5 — The system is tested on explanations it refuses to make

Asia FTD	Affiliate FTD	Affiliate p	Nearby event	Causal support
-8.13 pp	-15.81 pp	0.00463	Office Relocation	FALSE



Temporal proximity is context, not proof. FitzSight runs a falsification check before attribution.

Four trust boundaries make the Agent auditable

01 Local intent gate

Unsupported questions are refused before a model call.

02 Constrained planner

Only approved high-level actions; no SQL or arbitrary tool args.

03 Deterministic tools

Read-only SQL / Python own calculations, tests and decompositions.

04 Claim verifier

Evidence IDs, digest/status, ground-truth boundary and causal language.

Verification fails → FitzSight withholds the analytical answer.

Evaluation measures correctness, evidence quality — and safe refusal

Benchmark scenarios

5 / 5 PASS

Root-cause accuracy

100%

Evidence coverage

100%

Verifier violations

0

5 deterministic scenarios

CRM routing

Net deposits

Customer Intelligence

Lead quality

False correlation

8 adversarial cases / 8 PASS

Scope refusal

Planner policy

Missing evidence

Causal overclaim

_gt leakage

False correlation

The release gate rewards the system for refusing unsupported answers, not only producing correct ones.

Safety, compliance and reproducibility are product constraints

Operational boundaries

- Synthetic benchmark data only
- No real customer PII or employer data
- No investment advice or trading actions
- No automated AML / credit / suitability decisions
- Evaluation-only *_gt fields blocked from normal Agent SQL

Reproducibility

License

MIT

Backend

DuckDB

Fallback

SQLite

Planner

**Local
fallback**

Public repo

github.com/AplusNeutrino/FitzSight

BI tells you what changed. FitzSight investigates why the measurable evidence changed.

A reproducible financial investigation you can inspect, verify, challenge — and sometimes refuse.

Agent intents

5

Benchmark

5 / 5

Adversarial

8 / 8

Evidence coverage

100%

Evidence-grounded autonomous financial operations investigation