

Evidence-grounded Financial Operations Intelligence Agent

Brokerage / FinTech Operations Analyst

Autonomous investigation. Human decision.

GOAI 2026 · Boundless Agents · AI+ 金融




Not chat with a CSV.
A bounded investigation you can audit.

Dashboards show what changed. The investigation behind why is still manual.

Primary user

Brokerage / FinTech Operations Analyst



A generic LLM can write a plausible explanation quickly. Financial operations needs a reproducible evidence chain before narrative.



FitzSight converts a business question into a verified decision-support investigation



The model may select the next approved analytical branch. It never owns business arithmetic, arbitrary SQL, or high-impact financial actions.

Planner authority

**Approved actions
only**

Bounded adaptivity; local intent gate first

Calculation authority

**Read-only SQL /
Python**

Every competition-facing number is executable

Answer authority

EvidenceClaimVerifier

Verification failure → attribution withheld

Hero journey — evidence selects the next approved step

Question: “Why did European FTD conversion deteriorate after July 15?”

Actual execution

9 approved steps

Rendered from the verified v0.12 run JSON

Contribution drilldown

executed

Latency / anomaly scan

executed

Operational event check

executed

Document evidence

CRM-CHANGE-2026-0715#p1

Result-driven branching ≠ unrestricted autonomy. Every next step remains inside the approved action catalog.

FitzSight v0.12 — Verified Hero Run Evidence

Brokerage / FinTech Operations Analyst

Autonomous investigation. Human decision.

USER QUESTION

Why did European FTD conversion deteriorate after July 15?

VERIFIER: PASS | 6/6 claims verified

BOUNDED ADAPTIVE EXECUTION TRACE

P1 inspect_schema [executed]

Required operational fields were inspected.
Evidence: E0002

P2 query_affected_cohort [executed]

Affected Europe Team A+B cohort measured.
Evidence: E0003

P3 query_control_cohort [executed]

European control cohort measured.
Evidence: E0004

P4 statistical_validation [executed]

Conversion and response-time shifts tested.
Evidence: E0005, E0006

P5 contribution_decomposition [executed]

Team contribution decomposition completed; result controls the next approved branch.
Evidence: E0009, E0010

P6 anomaly_scan [executed]

Approved by contribution/statistical branch.
Evidence: E0011, E0012

P7 event_check [executed]

Operational event context checked after the latency/statistical signal.
Evidence: E0013

P8 document_evidence_check [executed]

Synthetic routing ticket paragraph linked by stable source reference.
Evidence: E0014

P9 evidence_boundary [executed]

Hero finding — evidence supports a CRM routing candidate, not a proven causal conclusion

Affected FTD

-7.53 pp

Europe Team A+B

Control

-1.21 pp

Other Europe teams

Response median

+29.15 min

Top contributor

Team A

contribution -1.87 pp

Root-cause status

supported

candidate · not causal proof

Evidence chain

Quantitative shift	Affected -7.53 pp vs control -1.21 pp
Latency signal	Median response +29.15 minutes
Contribution	Team A is largest negative contributor
Robust anomaly	8 post-change high-anomaly days
Operational context	Nearby CRM routing event
Source-addressable doc	CRM-CHANGE-2026-0715#p1

Evidence: E0013

P8 document_evidence_check [executed]

Synthetic routing ticket paragraph linked by stable source reference.

Evidence: E0014

P9 evidence_boundary [executed]

Fail-closed attribution boundary applied before verification.

Evidence: E0015

VERIFIED ANSWER

A supported root-cause candidate was identified.

- Affected Europe Team A+B FTD conversion changed by -7.53 percentage points after 2026-07-15.
 - Median response time in the affected cohort changed by 29.15 minutes.
 - The European control cohort conversion changed by -1.21 percentage points over the same split.
 - A CRM routing change is present in the nearby business-event log and is a supported root-cause candidate, not an automatically proven real-world causal conclusion.
 - Team A is the largest negative team-level contributor in the symmetric Europe FTD-rate decomposition (-1.87 pp contribution).
 - 8 of 27 post-change daily response-time medians exceed the robust pre-change high-anomaly threshold.
- Guardrail: The evidence supports the CRM routing change as the primary root-cause candidate in this synthetic benchmark; a real deployment should not convert temporal association into a causal claim without additional design/validation.

APPROVED FOLLOW-UP

What evidence supports the CRM routing change candidate?

The candidate is supported by the affected-vs-control FTD movement, statistical validation, response-latency signal, nearby operational event evidence, and synthetic document source CRM-CHANGE-2026-0715#p1. It remains a supported candidate rather than a proven real-world causal conclusion.

DOCUMENT EVIDENCE


CRM-CHANGE-2026-0715#p1 — Synthetic CRM routing change ticket

Effective 2026-07-15, Europe Team A and Team B move to the revised lead-routing queue. Operations should monitor assignment-to-first-response latency and FTD conversion during the post-change period.

Synthetic benchmark only. Product-process image is rendered from the actual verified runtime JSON; it is not a hand-authored analytical mock.

Failure branch — error evidence ends attribution, not the investigation

Tested dependency failure: the operational-event lookup raises an error Evidence record.



What fails

Operational-event dependency is unavailable. The system records the error and does not manufacture event or document corroboration.

What still succeeds

The evidence-backed metric findings remain valid; EvidenceClaimVerifier passes a guarded answer whose root_cause_status is insufficient_evidence.

Agent quality is not “always finding a cause.” It is completing the task with the correct uncertainty boundary.

Refusal story — temporal proximity is not evidence of causality

Question: “Asia conversion fell after July 20. Did the nearby office relocation cause it?”

Asia FTD

-8.13 pp

Affiliate FTD

-15.81 pp

Affiliate p

0.00463

Nearby event

**Office
Relocation**

Event causal support

FALSE

Tempting narrative

Office relocation

Asia FTD ↓

REJECTED

Measured driver

Affiliate quality ↓

Supported driver

false-correlation rejection = true

FitzSight is evaluated on explanations it refuses to make — not only the answers it produces.

Three additional workflows prove reuse — without diluting the hero narrative

01

Client-fund flows

Net-deposit / withdrawal concentration

Net change -\$187.8k

Observed driver only; no customer-motive inference

02

Customer Intelligence

Behavioral-value segmentation

Coverage 100%

Descriptive only; no AML / credit / eligibility decisions

03

Acquisition quality

Volume vs mix vs within-channel quality

Leads +315% · FTD -10.84 pp

Separates more leads from better leads

Main pitch = one complete investigation + one refusal. These workflows are reusable breadth, not five equal demo stories.

Technical depth comes from explicit authority boundaries, not more Agent labels

01 Local intent gate

Unsupported questions are refused before an external model call.

02 Constrained planner

Only approved high-level actions. No planner-generated SQL or arbitrary tool parameters.

03 Deterministic analytics

Read-only SQL / Python own KPI calculations, tests, decompositions and anomaly checks.

04 Source-addressable evidence

Tool outputs and synthetic operational documents enter one append-only Evidence Registry.

05 EvidenceClaimVerifier

Checks evidence integrity, _gt boundary and causal language before result delivery.

06 Human decision boundary

No trading, AML, credit, suitability or other high-impact financial decisions.

Verification failure is a product state: FitzSight may finish with a guarded answer or withhold the unsupported attribution.

Evaluation v2 asks whether the system generalizes — and whether the safety architecture matters

Fixed benchmark

5 / 5 PASS

Holdout runs

8 / 8Unseen seeds + question
paraphrases

Routing stability

100%

Evidence coverage

100%

Supported

75%1 CRM seed correctly
insufficient

Adversarial release gate

8 / 8 PASS

scope refusal

planner policy

evidence integrity

causal overclaim

_gt leak

false correlation

Controlled architecture ablation

Full FitzSight

100% adversarial refusal

No verifier/evidence gate

100% unsafe-answer rate**This is a controlled architecture ablation — not a Generic LLM baseline.**

One unseen CRM seed returning insufficient_evidence is retained as a robustness success, not hidden to make the metric look perfect.

Safety, traceability and open reuse are explicit product boundaries

Implemented today

- Synthetic benchmark data; no real customer PII
- Read-only analytical path with deterministic fallback
- Evidence Registry + source/paragraph document evidence
- Causal-language guardrail and fail-closed verifier
- No investment advice, trading or account actions
- No automated AML, credit or suitability decisions
- MIT license + public repo + tests + benchmark catalog

Production blueprint — not claimed as implemented

Identity / RBAC

Row- and field-level policy

PII masking

Retention / enterprise audit controls

Approved business-system connectors

Financial boundary: evidence-grounded decision support. The authorized analyst or institution retains the final judgment.

BI tells you what changed. FitzSight investigates what the evidence supports — and what it does not.

Autonomous investigation. Human decision.

Hero

CRM / FTD

bounded-adaptive investigation

Trust case

False
refuses unsupported causality
correlation

Holdout

8 / 8
routed + verified

Evidence coverage

100%
synthetic benchmark scope

A reproducible financial investigation you can inspect, verify, challenge — and sometimes refuse.